

KEY RESULTS DOCUMENT (KRD)

FE Loss Report Card

Loss visibility + dispute-before-debit for Pilots, in the Valmo Partner app

Fulfilment Product Space · F&E / Valmo / SX Claims & Losses · Doc version 7 · 3 Sep 2026 → v8

A Pilot who causes a shipment loss today finds out only when money leaves their payment — one line, no break-up, no reason, no way to respond. This KRD builds a **Losses** section in the Valmo Partner app that shows each loss with its proof, and lets the Pilot **accept it or dispute it before the money is cut**. Disputes are checked by human agents. Silence after fair notice means the amount is deducted.

Read Me — scope in one minute

- **Who it is for.** Pilots (also called FEs) only. Captains and other entities come later, on the same rails.
- **What it covers.** The in-app Losses section (three tabs: **Needs action / Debits / Wrong Pickups**), the accept/dispute flow, the human review via Kapture, and the data work to power it (debit-date column, day-0 Lost-in-Field feed).
- **What it does not cover.** Automating payout deduction for **SX-claims losses** — today executed manually via **file upload in the loss engine** — deferred; taken up after this ships. The attribution pipeline itself (stays as-is).
- **How this new solution differs from the previous one.** The earlier read-only two-card model (Actual vs Potential debit) is replaced by a three-way choice: accept / dispute / stay silent. This came from the leadership direction: show likely debits first, give the Pilot the right to respond, deduct only after.

Words used in this doc

Pilot / FE — the delivery partner using the Valmo Partner app. **AWB** — one shipment.

Gold table (SX Claims POD, gold.sx_report_debit) — where SX-claims losses live today, owned by the SX Claims POD.

Loss engine (Losses POD) — the central loss database (Loss Engine POD). Target home for all entity losses.

Payout engine — powers the Earnings/Payments tabs and executes deductions. Separate from the loss engine.

Kapture — the 3P ticketing tool where disputes go; 12–15 agents on a 3P payroll review them.

Co-Ownership at a Glance

Function	Owns	Bar
Product (PM)	Problem, solution, versioning, thresholds, post-launch validation	Pressure-tested logic; every edge case owned; risk register
Business / Biz-Fin	Selection % (Y), cool-off values, leakage budget, GTM, cost	Y/X/T set with data; leakage signed off
Design	Screens, states, comprehension	Pilot understands the 3 choices in 10–15s; ≥3 real Pilot walkthroughs
Analytics	Metrics, data contracts, dashboards, alerts	debit_date contract verified; funnel + overturn dashboards live
Engineering	Read/act services, Kapture integration, LiF day-0 feed, SLA clock	Auto-waive fires even if Kapture is stuck; no double-push

Function	Owns	Bar
Support / Ops (CX)	Agent rubric, queue, decision-tree reroute	Rubric + QC live before launch; no duplicate ticket paths

KR Snapshot

POD	F&E / Valmo / SX Claims & Losses
KR / Feature	FE Loss Report Card — loss visibility + dispute-before-debit
Current Version	V1 (first build) — draft, doc version 8
Status	KRD review → Kapture API + security discussion (PoC: Sanam Suri; Kapture: Priyanka)
Target timeline	To be updated
Contributors	PM [Shreyash] · Designer [TBD — V9 design under review] · Eng [TBD] · Analyst [Sanjay] · Business [Rachit] · Support [Nandini, Nikhil, Niharika]
Approvers	Product Mentor · EM · Design Mgr · Analytics Mgr · Biz-Fin · Biz Mentor · Security (Kapture image hand-off)
Linked docs	Flow diagram (FE_Settlement_Flow_v3) · Design V9 · Wireframes v3 · Business flow doc · Kapture call transcript (3 Sep '26)

SECTION I · Executive Summary

1. The Problem

Owner: Product; sized with Business / Biz-Fin.

Problem statement. When a loss is attributed to a Pilot, money is deducted from their payment with no explanation and no chance to respond. In the app they see one cumulative line (e.g. “Lost Shipments –₹255”) with no shipment-level break-up and no reason (Pilot app mock screens below). If the attribution is wrong, the Pilot has no way to contest it — the money is already gone.

What this causes.

- Pilots cannot correct the behaviour behind a loss they never see explained — so the same mistakes repeat.
- Every unexplained deduction becomes a support query. Deduction break-ups are pulled and explained manually by the support team today (support team, Kapture call 3 Sep '26).
- Deducting before hearing the Pilot means wrong attributions cost Pilots money with no recourse — a fairness and trust problem, and a grievance/attrition risk.

Size. ~₹800 debited per Pilot per month across the ~4K Pareto Pilots (372 hubs) → ~₹32L/month, ~₹3.8Cr annualised, before attribution coverage extends to the full base [analytics, Jun '26]. ~3 attributed loss events per Pilot per month on average [analytics, Jun '26]. Full-base sizing [TBD].

Mock example of currently live Pilot Payment section in Valmo1/Valmo Partner application

Deductions section not showing any AWB level details	Adjustment section not showing any AWB level details
< Payment Details Help Payment for 19 May ₹234 Base Pay ₹3200 Delivery 240 Deliveries @ ₹13/Delivery ₹3120 100 Deliveries @ ₹14/Delivery ₹1400 10 Consolidated Deliveries @ ₹10/Delivery ₹100 Pickup 240 Pickup @ ₹13/Pickup ₹3120 10 Additional Pickups @ ₹10/Pickup ₹100 Incentives + ₹37 Adjustments + ₹50 Deductions - ₹120 TDS 1% - ₹48 Facilitation Fee - ₹48 Lost Shipments - ₹255	< Payment Details Help Payment for 19 May ₹234 Base Pay ₹3200 Delivery 240 Deliveries @ ₹13/Delivery ₹3120 100 Deliveries @ ₹14/Delivery ₹1400 10 Consolidated Deliveries @ ₹10/Delivery ₹100 Pickup 240 Pickup @ ₹13/Pickup ₹3120 10 Additional Pickups @ ₹10/Pickup ₹100 Incentives + ₹37 Adjustments + ₹50 Missing Shipment Credit ₹20 Previous Day Correction ₹20 Discrepancy Adjustment ₹10 Deductions - ₹120

2. High-Level Solutioning

Owner: Product.

Direction. “Innocent until proven guilty” (leadership direction, Aug '26): show the Pilot the likely debit first, give them the right to respond, and deduct only after — by their acceptance, by a failed dispute, or by silence after fair notice.

Hypothesis. If a Pilot can see each loss with its proof and act on it (accept or dispute) before money is cut, then wrong attributions get corrected and repeat mistakes fall, thereby reducing loss events per active Pilot. **The flywheel:** visibility → Pilots fix behaviour and contest wrong attributions → SX claims go down → SX improves → CPS reduced

The model in summary: A loss is attributed and shown in the app with the Pilot's own photos. The Pilot has **7 days from the day it is attributed to them** (one clock for every case; no separate seen-based window) to choose:


If they choose **Accept** (“yes, my mistake” + a reason chip — business then selects **Y% of accepted shipments** for full deduction; the rest are not deducted; the Pilot is told the outcome before payout, and the app never promises a waiver),

If they choose **Dispute** (one chance, reason chip + optional text note — **no voice notes** — goes to Kapture, a human agent reviews 100% of disputes within a 7-day SLA; right → marked loss waived off; wrong → full loss marked value deducted; **SLA missed** → **auto-waived**), or **stay silent** → full amount deducted at window close.

A Pilot whose last **X disputes in a row** were found wrong gets disputes paused for **[T] days** (accept and silence stay open; one won dispute resets the counter).

The three tabs.

- **Needs action** — open losses, sorted by days left; act here.
- **Debits** — what was deducted and why, grouped by date, each item pointing to the exact line in the payment.
- **Wrong Pickups** — information only; never debited.



Lost in Field (LiF) is special. A parcel taken out but neither delivered nor returned is auto-deducted the **same day** by the payout engine, and the Pilot has 7 days to return it to the hub to get the money back. So it appears in **both tabs from day 0**: a recovery card in Needs action (“₹40 already cut — return by <date> to get it back”) and a live row in Debits. On return, the card leaves Needs action and a green credit row is added in Debits — like a bank statement.

Scope / phasing.

- **V1** — the three tabs; the accept / dispute / silence flow; Kapture ticket integration (two-way); Lost-in-Field day-0 in both tabs; SMS at attribution + day-4/6 reminders; debit-date column in the loss engine; simple-English copy, localised to the app's language set. **Also in V1 (v8):** deep-link from a payment's deduction line to the matching loss, at line-item level (F27).
- **V2** — central loss DB in the loss engine as the single read path for all entity losses; Captain report card on the same rails. Push notifications for needs action/debits/wrong RVP..

- **Deferred** — automating payout deduction for **SX-claims losses** (today a manual file upload in the loss engine) — after this ships. **In design:** fraud-awareness overlay stays — treatment to be refined with the designer.

3. Goals, NSM — Success & Check Metrics

Owner: Product; computation with Analytics.

Element	Definition
KR goal	Reduce losses from seller claims, fairly: fewer repeat mistakes (behaviour change) and wrong attributions corrected before money is deducted.
NSM	Loss (claims) events per active Pilot per week , exposed vs holdout, trending down. Robustness cut: loss events per 50 reverse-pickup shipments. Clarification: the NSM counts debited loss types only (photo junk, mismatch, Lost in Field). Wrong pickups carry no deduction (only tracked for now) separately as a wrong-pickup rate (advisory), outside the NSM. The per-50-RVP cut normalises by reverse-pickup volume (exposure), not by wrong-RVP events.
Model-health metrics	Action rate (% of marked losses acted on in-window); accept rate; dispute rate; silence rate; overtake rate (disputes won) by reason and by agent; SLA compliance (% disputes closed in 7 days) and auto-waive count; Lost-in-Field recovery rate (% returned in window).
Leading indicators	Loss-bearing-Pilot view rate; view-to-improvement delta (defect-rate drop in the 3–7 days after a view).
Money metrics	Total loss events and ₹, exposed vs holdout; selection leakage (₹ not deducted on accepted-but-not-selected shipments) — a budgeted cost, signed by Biz-Fin [TBD].
Expected impact	[TBD] once baseline locks
Non-goals	No payout automation; no appeal-on-appeal; the old Potential-debit concept is gone; other debit entities out (own KRDs);awareness overlay — design treatment with the designer [open].

Guardrails (each wired to a threshold, window, action, owner).

Guardrail	Breach	Window	Action	Owner
Wrong-Pilot data / photo shown (privacy)	Any confirmed case	Immediate	Alert → pause tab (feature flag) → root-cause	PM + Eng
Deducted amount ≠ what the app showed (reconciliation)	Any confirmed mismatch	Weekly	Alert → reconcile → pause if systemic	Biz + Eng+PM
Dispute volume > agent capacity (SLA at risk)	[TBD]% queue breach	Daily	Add capacity / tighten dedupe; auto-waive still fires	Support + Biz
Grievance / escalation spike	[TBD]% over holdout	2 weeks	Investigate → pause if sustained	Ops + Biz
App performance (low-end / 2G)	> NFR target	Continuous	Fix → pause at scale	Eng

Kill-switches. Feature flag on the Losses section (by cohort/DC). The Kapture ticket feed and the LiF day-0 feed each have their own off-switch — turning one off never blocks the auto-waive clock (it runs in our system).

4. Opportunities to Use AI [Contextual]

Owner: Product.

None in V1 (show, choose, review). V2 candidates: agent-assist on dispute review (photo checks), repeat-pattern coaching. Identify patterns of fraudulent pilots to refine the debit process further.

SECTION II · Product Requirement Documentation

5. User Stories and Capabilities

Owner: Product.

5a. User Stories

ID	As a...	I want to...	So that...	Pain today
US-01	Pilot	see every loss marked to me, with my own photos and the reason	I know what went wrong and how much is at stake	One cumulative deduction line; no break-up, no reason
US-02	Pilot	dispute a loss I did not cause, before money is cut	a wrong attribution does not cost me	Money is cut first; no way to respond
US-03	Pilot	accept a genuine mistake with my reason	the case closes and I know the outcome before my payout	No channel to respond at all
US-04	Pilot	see a lost parcel the day (or D+1) it is marked, with the return deadline	I can return it and get the money back	I learn about it days later, after the window is gone
US-05	Pilot	match each deduction in my payment to a shipment and reason	my payment makes sense without calling support or captain	Support pulls the break-up manually
US-06	Support agent	receive disputes as structured tickets with evidence attached	I can decide cases consistently and fast	Queries arrive as free-form WhatsApp/support messages
US-07	Pilot	tap a deduction line in my payment and land directly on that exact loss	I never have to hunt for why money was cut	The payment line and the loss explanation are unconnected screens today

5b. Key Functional & Non-Functional Requirements

Owner: Product. P0 = V1-blocking, P1 = fast-follow.

Surface & tabs

#	Requirement	Pri	Notes / edge cases
F1	Resolve the Pilot from the authenticated app session (stable ID; fallback = phone + name where unique ID is absent)	P0	Unmapped ID → clean zero state, never another Pilot's data
F2	Losses appears as the third segment in the existing Earnings screen (My Earnings / Payments / Losses); rate card hidden while Losses is active	P0	Badge on the segment when open cases exist

#	Requirement	Pri	Notes / edge cases
F3	Three tabs, in this order: Needs action / Debits / Wrong Pickups	P0	Tab names are config (vernacular pass §11)

Needs action (open cases + window)

#	Requirement	Pri	Notes / edge cases
F4	Show each open loss with reason (plain words), ₹, AWB, and days left. Sort by days left, ascending	P0	Amount never reorders; new items land by their own deadline
F5	Window = 7 days from the day of attribution — one clock for every case. Window state drives the countdown chip	P0	First-view event still logged (F26) — analytics only; it does not drive the window (v8)
F6	Badge and “at-risk” total count actionable disputable items only — exclude decision-pending cases and Lost-in-Field recovery cards	P0	LiF money is recoverable, not at risk [Open decision — v8] exact counting rules to be finalised (see 6h).
F7	Acted-on cases stay visible with a Decision pending chip and a 3-step tracker (raised → your reason → decision expected). A Closed strip at the bottom keeps no-money outcomes (dispute won, accepted-not-selected) visible	P0	Tracker for accepted cases says “decision before your payout” — no dated reply promise

Loss detail (evidence first)

#	Requirement	Pri	Notes / edge cases
F8	Detail shows the Pilot’s own photos above the choices; image set per reason-code mapping (mismatch = 4: 2 pickup + 2 QC; photo-junk = 2; Lost-in-Field = none)	P0	Image missing → placeholder + reason + tip; row never blocked
F9	Two option cards, each stating what can happen: Dispute shows both endings (₹0 if right / full ₹ if wrong); Accept states the process + confirmation promise. Silence consequence printed on-screen	P0	Copy never promises a waiver, never shows odds or

#	Requirement	Pri	Notes / edge cases
			Y% [Open decision — v8] how the silence consequence is displayed on-screen (see 6h).
F10	Voice-out button on every explanation block (text-to-speech in the app language)	P0	Low-literacy support

Accept path

#	Requirement	Pri	Notes / edge cases
F11	Accept = one mandatory reason chip (structured) + optional note; “Other” requires a note. Case closes on confirm	P0	Chips feed coaching segmentation
F12	Business selection runs on accepted cases: Y% of accepted shipments (not Y% of amount) are selected for full deduction; the rest are not deducted. Y and the selection logic are business-internal config	P0	Identical acceptances can end differently — FAQ owns this question (§14, §22)
F13	Proactive confirmation before payout: whichever way selection lands, tell the Pilot in-app + SMS before the payout adjusts	[Open decision]	Proposed by Product; firmness, channel (in-app vs SMS) and timing to be decided (v8, see 6h)

Dispute path (Kapture)

#	Requirement	Pri	Notes / edge cases
F14	Dispute is single and final per loss: one mandatory reason chip + optional text note — no voice notes . No photo upload — the parcel is gone and the evidence is already on screen	P0	“Other” requires a note
F15	On submit, create a Kapture ticket (AWB-tagged, evidence links, Pilot’s reason); agent decisions flow back via the update-ticket API and update the case + notify the Pilot	P0	One ticket per AWB; if multi-select ever ships, parent ticket + per-AWB child cases

#	Requirement	Pri	Notes / edge cases
F16	The 7-day dispute SLA clock runs in our system, not Kapture's. On breach, the case auto-waives (₹0) regardless of ticket state	P0	A stuck 3P queue must never become a stuck debit
F17	Cool-off: > X consecutive disputes found wrong → dispute option paused for [T] days (shown with reason + reopen date); accept and silence stay open; one won dispute resets the counter to 0	P0	X, T = business config [TBD]

Lost in Field (day-0, both tabs)

#	Requirement	Pri	Notes / edge cases
F18	Ingest a day-0 (or by D+1) Lost-in-Field signal from the payout engine (side-channel until the loss engine holds provisional records). Show a recovery card in Needs action (“₹X already cut — return to your hub by <date> to get it back”) and a live row in Debits	P0	SDE aligned (Sep '26); exit = hub return
F19	On return in window: remove the card from Needs action; keep the Debits debit row (status “Returned ✓”) and add a green credit row in the next cycle — bank-statement style, nothing deleted	P0	Credit event must arrive on the same channel
F20	Shared dedupe key (AWB + debit event) so the day-7 loss-engine record merges into the same case — never a duplicate row	P0	Dual-source integrity, NFR-7

Debits, Wrong Pickups, notifications

#	Requirement	Pri	Notes / edge cases
F21	Debits grouped by date (business call, Sep '26), newest first; every row states its path (accepted / dispute checked / no action / auto-deducted) and sits under its debit date	P0	Debit date = the payout engine's execution date — see the data contract in 6b
F22	Every debited item carries a payment pointer — the exact payout date + section + line it appears in (claims → Adjustments · “Shipment Loss for Junk/Mismatch”; LiF → Deductions · “Lost Shipments”). Pointer text is real data, not hardcoded copy; mappings are config	P0	Line labels must match the payment screen word-for-word
F23	History kept 3 months in-app (data persists in the loss engine beyond)	P0	Older → scroll; no calendar picker
F24	Wrong Pickups tab: information only, never debited; shows the Pilot's recorded responses + optional “Add your side”; reverse-chronological; no red, no minus signs	P0	Repeated cases may be reviewed by the hub — stated honestly
F25	Notifications — P1, feasibility itself to be decided. Proposed: SMS at attribution + day-4/6 reminders; push also P1→V2. Whether SMS/notifications are possible is an open decision ; working	P1 [decision]	Window validity no longer depends on

#	Requirement	Pri	Notes / edge cases
	assumption meanwhile: Pilots open the app regularly since payments are visible only there. Day-0 Lost-in-Field SMS ask: keep or drop [open] (owner: opstech/payout).		notifications — 7 days from attribution (v8)
F26	Emit events: case viewed (analytics only) option opened, accept/dispute submitted, decision received, credit posted	P0	Powers §3 metrics + §7
F27	Deep-link (V1 — v8) . Tapping a loss deduction line in Payment Details opens the matching case in the Losses tab, at line-item level. Reverse of the F22 pointer.	P0	Payments-surface change — payout front-end team dependency (6f)

Non-functional requirements (each a pass/fail UAT).

- **NFR-1 (low-end device)**: section renders < [TBD]s on a 2GB-RAM Android; text renders before images.
- **NFR-2 (2G)**: thumbnails lazy-load; full image on tap; summary and lists usable before images resolve.
- **NFR-3 (offline)**: last-fetched view cached and readable with an “as of <date>” state [frontend feasibility check pending].
- **NFR-4 (image serving, two rules)**: in the Pilot app, images use the public links already stored in the gold table (decision, Jul '26). Into **Kapture (3P)**, images go via a **broker service issuing expiring links (2–4h)** — never raw links to a third party [security sign-off pending]. **[Marked open: tracked in 6h so it is not forgotten.]**
- **NFR-5 (reason scalability)**: reasons rendered are controlled by a whitelisted-reasons array (config); the list may grow to ~10 without redesign.
- **NFR-6 (SLA correctness)**: the auto-waive fires exactly once per breached case, even if Kapture is unreachable; a late agent decision after auto-waive cannot un-waive.
- **NFR-7 (dual-source integrity)**: while LiF arrives via the side-channel (payouts engine), merges with loss-engine rows must never double-count (F20 key) and totals must still match the payment lines to the rupee.

6. Solution Deep Dive

Owner: Product (validated with Eng, Analytics, Support).

6a. Solution Pathways — Trade-offs

Option	Pros	Cons
A. Read-only transparency (v6 model)	Simple; no ops load	Pilot still has no voice; wrong attributions still cost money; grievance risk stays
B. Dispute-before-debit with accept option (chosen)	Fair process; wrong attributions corrected before money moves; disputes audit the attribution pipeline for free	Agent workload; SLA machinery; selection fairness must be communicated well
C. Sampled merit review (X% of	Caps agent load by design	Considered and dropped: with the accept option, dispute volume

Option	Pros	Cons
disputes checked, rest auto-waived)		self-sorts to genuine cases, so 100% review is affordable and cleaner

6b. Solution Details

Current state (unchanged by this KRD). Hub-checkpoint images → analyst audit → audited data → gold.sx_report_debit assigns the debit entity, reason and cycle date → Finance executes deductions. Lost-in-Field is separate: detected and auto-deducted the same day by the payout engine; the loss engine receives only unresolved cases, 7 days later; returned-in-window cases are auto-credited and never recorded.

Core rules of the model.

- **Every loss gets a hearing window. 7 days from the day of attribution.** Notification approach is an open decision (F25); the working assumption is Pilots open the app for payments. Silence after that = full deduction — a real choice, not a trap.
- **Accept = settlement.** One reason chip; case closes. Business selects **Y% of accepted shipments** for full deduction; the rest are not deducted. The app promises the process and the date, never the outcome — no odds, no Y% on screen. The Pilot is told the result before payout (F13). **[Confirmation approach — open decision]**
- **Dispute = one fair shot.** 100% human review in Kapture, 7-day SLA, both endings stated before submit. Right → ₹0. Wrong → full amount. **SLA missed → auto-waived** — our delay is never the Pilot's cost.
- **Cool-off protects the queue.** Only **consecutive** wrong disputes trigger it; a single win resets it; accept stays open throughout. It throttles gambling, never removes recourse permanently.
- **Lost in Field = recover, not argue.** Already deducted upstream; the app's job is to surface the 7-day return remedy on day 0 and to pair the debit with its credit when the parcel comes back.
- **Wrong pickups are advice, not debits.** Shown with the Pilot's own recorded responses; never money.
- **Status lifecycle (data).** ATTRIBUTED → (viewed) → IN-DISPUTE / ACCEPTED → DEBITED / WAIVED / NOT-DEDUCTED, plus LiF RETURNED-CREDITED. This lifecycle is the shape the central loss DB should hold for every entity — Lost-in-Field day-0 is just its first provisional writer.

The reconciliation contract (load-bearing).

- A new **debit_date** column is added to the loss engine (this KRD). It must equal **the payout engine's execution date** — the "Payment for <date>" date, not the credited date, not an ingestion timestamp.
- Debits post on their **resolution date** (a loss attributed 7 Aug, resolved 11 Aug, shows under 11 Aug). The detail screen's timeline carries the attribution date, so nothing double-counts.
- Claims losses reach the payment via the adhoc-adjustment rail → they appear inside **Adjustments** as "Shipment Loss for Junk/Mismatch" (a cumulative negative line per the Pilot's payout period). LiF appears inside **Deductions** as "Lost Shipments". The app's per-item payment pointer (F22) must name these exactly. Ask logged: move claims losses to Deductions later; **line-level match is the contract either way**.

Edge cases.

- Pilot never opens the app → the 7-day window runs from attribution and closes on schedule; deduction proceeds. **[Open decision]** whether SMS/notifications are feasible; working assumption: **Pilots open the app regularly since payments are visible only there; the case remains explained in Debits.**
- Dispute submitted at day 6, resolved day 13 → posts under its resolution date; tracker shows the journey.
- Kapture down / ticket stuck → SLA clock (ours) still fires auto-waive; ticket reconciled later (NFR-6).

- Agent decision arrives after auto-waive → waive stands; decision logged for QC only.
- LiF returned but hub scan failed → case lands as “not returned”; the Debits detail carries “Returned already? Raise a ticket” — the escape hatch for exactly this error.
- Same AWB from side-channel and loss engine → F20 key merges; never two rows.
- Reason code missing from tip/image mapping → generic tip / no-image state; gap logged.
- BTS / peak traffic: image serving degrades to text-first (NFR-1/2); lists paginate. **[Open decision]** exact pagination/degradation behaviour at peak (see 6h).

6c. Operating-Model Coverage Matrix

Support operating model changes: dispute tickets shift from free-form WhatsApp decision-tree entries to structured in-app tickets. The decision tree must **reroute** its payment-dispute path to the app (owners: Nikhil, Niharika) or double-ticketing breaks the SLA/auto-waive logic. Agent rubric + QC = new ops artefacts (6k).

Ticket-channel map — where this feature does and does not interact:

Channel	What it carries today	Interaction with this feature	Change needed
In-app Losses tab (new)	—	The only channel for disputing an attributed loss	New (this KRD)
WhatsApp decision tree	Payment/deduction queries land here	Must NOT accept loss disputes after launch — double-ticketing breaks the SLA/auto-waive clock	Reroute payment-dispute path to the app — launch gate
Email-based tickets	General support / escalations [to verify with Support]	No loss-dispute handling; agents redirect to the app	Redirect macro + agent brief [TBD]
Helpline / captain escalation	Ad-hoc queries on the floor	Explain + redirect to the app; captains briefed (§14)	Briefing content only

6d. Designs

Owner: Design. Reference build: FE Report Card V9 (by PM); designer review pending.

Screens covered in V9: Needs action (in-situ, sorted); loss detail with two option cards; accept sheet + accepted state; dispute sheet; decision-pending tracker; dispute-won; Debits ledger with credit rows; closed-case receipts; cool-off; Wrong Pickups with catalog comparison; Lost-in-Field; empty state; voice-out on all explanation blocks. Copy corrections in flight from KRD review: tracker’s accepted-case promise (“before your payout”, not a dated reply), dispute card’s two endings restored, tab order (Debits second). Design bar: a Pilot understands the three choices in 10–15s; ≥3 real Pilot walkthroughs before sign-off.

<https://claude.ai/design/p/d575b3af-5b11-48a6-807e-1b697967139f?file=FE+Report+Card+V9.dc.html&via=share>

6e. Tech Feasibility Checks

Point	Concern	Implication	Status
LiF day-0 feed from payouts engine	Case lives in the payout engine on day 0; loss engine gets it day 7+	Side-channel read + F20 dedupe key + credit event on the same channel	SDE thinking through (Sep '26)
Kapture APIs	Create-ticket + update-ticket; Meesho instances are customised per team	Contract discussion with Kapture (Priyanka) once this KRD is shared; two-way sync	Open
Image hand-off to 3P	Raw links to a 3P are not acceptable	Broker service, expiring links (2–4h); security sign-off	Open (Sanam to loop security)
SLA clock ownership	Breach → auto-waive must fire even if Kapture is stuck	Clock + waive trigger in our system; ticket state reconciled after (NFR-6)	Open
debit_date column	New column; must equal the payout execution date	One source event for both systems; UAT-verified line-level equality	Open (Analytics)
Status lifecycle in the central DB	Loss engine today stores settled records only	Provisional records + status transitions — same shape the dispute flow needs anyway	Open (Loss Engine POD)

6f. Dependencies & Cross-POD Alignment

Owner: Program + Product. "Told" ≠ done — verify delivered.

Dependency	Owner	Needed by	Status
debit_date column in the loss engine (= payout execution date)	Analytics + Loss Engine POD	Build start	[TBD] / No
Day-0 LiF signal + return/credit event + dedupe key (F18–F20)	Payout engine + Eng	Build start	SDE apprised / No
Kapture create/update ticket APIs + instance customisation review	Kapture (Priyanka) + Sanam	Build start	Call done 3 Sep / No
Image broker with expiring links + security sign-off	Eng + Security (via Sanam)	Before UAT	[TBD] / No
Support decision-tree reroute (payment-dispute path → app)	Nikhil + Niharika (Support)	Launch	[TBD] / No
Agent rubric + QC + queue setup (12–15 3P agents, existing capacity)	Nandini (Support) + Product + Biz	Before UAT	[TBD] / No
Day-0 LiF SMS from the opstech deduction flow	Opstech / payout team	Launch	Ask logged / No → Open (v8): keep or

Dependency	Owner	Needed by	Status
			drop this ask [decision]
Reason-code → tip and → image-set mappings (config)	Business	Build start	[TBD] / No
Y%, X, [T] values + selection-leakage budget	Business	Before experiment	[TBD] / No
Payment line labels confirmed (“Shipment Loss for Junk/Mismatch” in Adjustments; “Lost Shipments” in Deductions); section-move ask logged	Payout team + Finance	Build start	Confirmed Sep '26 / ask open
Payment Details → Losses deep-link at line-item level (F27)	Payout front-end team + Eng	Build start	[TBD] / No

6g. Implications on Ground Operations

No new capture step — the delivery/pickup flow is unchanged. Ground implications are behavioural:

- (1) Pilots must understand three choices, in vernacular, fast — FGD-tested;
- (2) identical acceptances can end differently (selection) — floor comparisons will happen; captains and the FAQ must own the answer (“every accepted case goes through the same review”);
- (3) LiF recovery depends on the hub return loop working — captains briefed to accept and scan returns promptly..

6h. Key Alignment Points

Decided.

Call	What was decided	Decided with
Dispute SLA breach	Auto-waive — our delay never costs the Pilot	Anant + team, Aug '26
Selection on accepted cases	Y% of accepted shipments , full deduction on selected; never Y% of amount; outcome confirmed before payout	Business, Sep '26
Silence	Full deduction at window close — with SMS + reminders, seen-based window, 14-day cap as protections	Business, Aug '26
Cool-off	> X consecutive wrong disputes → pause [T] days; win resets; accept stays open	Team, Aug '26
Lost in Field	Day-0 in both tabs; hub return is the exit; on return — leave Needs action, bank-statement credit pair in Debits	Team, Sep '26
Debits grouping	By dates (not payout periods) — business call; item-level payment pointer carries reconciliation	Business, Sep '26

Call	What was decided	Decided with
Dispute inputs	One dispute, chips + text/voice; no photo upload (parcel gone; evidence on screen)	Team, Sep '26
Posting basis	Resolution date (closes the old claim-date vs debit-date question)	Team, Sep '26
Dispute window basis	7 days from attribution for every case — no seen-based clock, no 14-day cap	Product + Business, Sep '26
Payment deep-link	Line-item deep-link from Payment Details into the Losses tab ships in V1 (F27); push notifications move to V2	Product, Sep '26
Dispute inputs	Chips + optional text note only — no voice notes	Product, Sep '26
Reason naming	“Delivery photo not clear” is ICUD junk — same reason, FE-facing name	Product + Business, Sep '26
Awareness overlay	Stays in the design; copy/treatment refined with the designer	Product + Business, Sep '26

Open.

- Y%, X, [T] values and the leakage budget [Business + Biz-Fin].
- Fraud-awareness overlay — in or out, copy, and triggers [debate scheduled; concerns logged: tone on the losses surface, force-show after a lost dispute].
- Kapture instance customisation + API contract details [Priyanka].
- Security sign-off on the image broker [Security via Sanam].
- Offline caching feasibility [Frontend].
- **Notification feasibility**: whether SMS/push are possible at all; if yes, templates + cadence; day-0 LiF SMS — keep or drop [Product + Eng + opstech].
- **Pre-payout confirmation, F13**: firm requirement or not; channel and timing [Product + Business].
- **Badge / at-risk counting rules, F6**: exact inclusion rules to finalise [Product + Design].
- **Silence-consequence display, F9**: how it is shown on-screen [Product + Design].
- **BTS / peak degradation behaviour**: pagination + image fallback specifics [Eng].

6i. Versioning Calls

When	Item	Why
V1	Full model as scoped in §2 + payment deep-link	The choice flow and LiF day-0 are the product; shipping read-only first would rebuild trust twice
V2	Central-DB single read path; Captain report card + push notifications	Same lifecycle, next entities
Deferred	Automating payout deduction for SX-claims losses — today a manual file upload in the loss engine	After the report card is done (team call, Jul '26)
In design	Fraud-awareness overlay	copy/treatment refined with the designer

6j. Pressure-Testing Thresholds & Configuration

High-blast-radius config, all [TBD], none shipped from gut feel

Y% (selection rate — sets both deterrence and leakage; sized with Sanjay on real claim distributions),

X (consecutive wrong disputes — too low punishes unlucky honest Pilots),

[T] (cool-off length — one debit cycle is the starting suggestion), window length (**7 days from attribution**), reminder days, whitelisted reasons, tip/image mappings, payment line labels. All config, no hardcoding.

6k. Manual Interventions

Dispute review is deliberately manual: 12–15 3P agents on the existing queue, working to a written **rubric** (reason required on both accept and reject decisions; correct-loss vs wrong-loss classification) with **maker/checker QC** — with auto-waive on SLA, agent throughput converts directly to waived money, so per-ticket contract incentives must be checked. Mappings (tips, image sets, labels) are maintained lists with named owners. The support decision tree needs upkeep after the reroute.

6l. Risk Register

#	Risk	Likeli.	Impact	Mitigation	Owner
R1	Pilot never learns of the window → silent lapse feels like today's system with extra steps	Med	High	Notification feasibility is itself open (F25, v8) — until decided, discovery rests on in-app visits (payments live there), the segment badge and the bell. If notifications are ruled out, revisit window length. Deep-link from payments (F27) adds a discovery path in V1.	Product + Eng
R2	Wrong-Pilot data or photo shown (privacy)	Low	Very High	Authenticated ID join; immediate-pause guardrail	Eng + PM
R3	Dispute flood → SLA breaches → auto-waive leaks money	Med	High	Accept option self-sorts volume; evidence-first UX; cool-off; capacity guardrail (§3); queue dashboards	Support + PM
R4	Agent inconsistency → identical cases decided differently	Med	High	Rubric + reasons both ways + maker/checker QC + overturn-by-agent metric	Support + Product
R5	Selection fairness — identical acceptances end differently → floor anger	Med	Med	Process-language copy; FAQ + captain briefing own the answer; proactive confirmation (F13)	Product + Biz
R6	Double-ticketing via the old WhatsApp path → two SLAs, broken clocks	High	Med	Decision-tree reroute is a launch gate (6f)	Support
R7	LiF dual-source drift → duplicate rows or totals ≠ payment	Med	Med	F20 dedupe key; NFR-7; line-level UAT	Eng + Analytics

#	Risk	Likeli.	Impact	Mitigation	Owner
R8	Kapture dependency (API custom work, security) slips the timeline	Med	Med	Contract talks start with this KRD; broker scoped early; fallback = manual queue bridge for pilot cohort	PM + Biz + Kapture
R9	Low adoption (pull-only surface) → windows lapse unseen	Med	Med	SMS carries the surface; badge + bell; leading-indicator tracking from day one	Product

7. Data Instrumentation

Owner: Analytics. [pending]

Events: loss_case_viewed (starts the window), option_card_opened, accept_submitted (chip), dispute_submitted (chip), kapture_ticket_created / decision_received, auto_waive_fired, selection_result_shown, lif_recovery_shown / returned / credit_posted, payment_pointer_tapped — each with pilot_id, awb, reason, ₹, timestamps.

Funnels: marked → viewed → acted → outcome; dispute → decision → overturn.

Dashboards: §3 model-health set + SLA/queue + reconciliation counter. GATE: events reviewed before dev.

8. Cost Implications [Contextual]

Owner: Engineering + Biz-Fin. [pending]

Agents: existing 12–15 3P capacity, incremental utilisation [TBD with Support]. Kapture: per-ticket / customisation cost [TBD via Sanam]. One-time: broker service, LiF side-channel, app build. Running: SMS (attribution + 2 reminders + confirmations) at [TBD]/case; selection leakage is a policy budget (§3), not a system cost.

SECTION III · Launch Readiness

9. Experimentation & Impact Measurement

Owner: Product, with Analytics. Tooling [TBD].

Element	Detail
What we're testing	Does see-and-respond reduce loss events per Pilot vs no visibility? Causal question → test-control: Losses section ON for exposed Pilots, OFF for holdout.
Who's in the test	Unit: Pilot × week. Random split within similar groups (tenure, Pareto DC, baseline loss rate, reverse-pickup volume).
Main metric (NSM)	Loss events per active Pilot per week, exposed vs holdout. Robustness: per 50 RVP shipments.
Early signal	View-to-improvement delta — defect-rate drop in the 3–7 days after a case view.
Also read	Action rate; accept vs dispute vs silence mix; overturn rate by reason (an audit of the attribution pipeline itself); SLA compliance; LiF recovery rate; leakage ₹.
Guardrails	§3 table — privacy, reconciliation, queue capacity, grievance, app performance.
How big, how long	[TBD] with Analytics (power calc). Keep a long-term holdout — confirm the change sticks, not a novelty dip.
Ship rule	Scale only if the NSM is down vs holdout at [TBD] significance with guardrails clean; otherwise iterate or hold. Test-control is mandatory.

10. User Acceptance Testing [Pre-Prod]

Owner: Product & Design. [pending]

#	Scenario	Expected	Result
1	Window starts on first view	Countdown begins at attribution for every case; window closes at day 7 regardless of whether the case was viewed	TBD
2	Sorting + badge	Needs action sorted by days-left ascending; badge counts actionable items only (pending + LiF excluded)	TBD
3	Evidence-first detail	Photos render per reason mapping (4 / 2 / none); missing image → placeholder, row usable	TBD
4	Accept end-to-end	Chip required ("Other" needs note); case closes; selection outcome communicated per the decided confirmation approach (F13 — open)	TBD
5	Dispute end-to-end	One dispute only; Kapture ticket created with evidence links; agent decision returns and updates the case + notifies	TBD

#	Scenario	Expected	Result
6	SLA breach	Unresolved at day 7 → auto-waive fires once; late agent decision cannot un-waive	TBD
7	Cool-off	Xth consecutive wrong dispute pauses disputes with reason + reopen date; accept still works; a win resets the counter	TBD
8	Silence	No action by window close → full deduction; Debits row reads “no action taken”	TBD
9	LiF day-0	Recovery card + Debits row on day 0; excluded from at-risk total; day-7 loss-engine record merges (no duplicate)	TBD
10	LiF return	Card leaves Needs action; Debits keeps debit row (Returned ✓) + green credit row next cycle	TBD
11	Reconciliation	For sampled payouts: sum of app items per payment line = the payment line, to the rupee, labels word-for-word	TBD
12	Decision-tree reroute	Old WhatsApp payment-dispute path redirects to the app; no duplicate ticket for the same AWB	TBD
13	Wrong Pickups	Info-only; no ₹ minus anywhere; “Add your side” saves	TBD
14	Notifications	Per the decided notification approach (open): if SMS enabled — attribution + day-4/6, app-language copy, deep-open (v8)	TBD
15	Offline / 2G / low-end	Cached “as of” view; text before images; NFR-1 timing	TBD
16	Privacy	Unmapped ID → zero state; never another Pilot’s case or photo; Kapture links expire	TBD
17	Voice-out + languages	Every explanation block reads aloud in the app language; fallback to default	TBD
18	Design fidelity	Matches approved V9-derived Figma; ≥3 Pilot walkthroughs done	TBD
19	Payment deep-link (v8)	Tap a loss deduction line in Payment Details	Opens the Losses tab at the matching case (line-item level); back returns to the payment

Exit: zero open P0/P1 defects.

11. Feature Training & Content

Owner: Product (content) + Training. [pending]

Simple-English base copy → vernacular pass for every app language (tab names, the three choices, LiF recovery line). Comprehension check: three choices + “what happens if I do nothing” understood in 10–15s. Agent-side: rubric + decision-reason templates.

12. Demo Day [Pre-experiment in Prod]

Owner: Product & Design. [pending]

Real Pilots, real cases: watch the choice moment, the accepted-case confirmation, and one LiF recovery. Feed back before go-live.

13. Data Instrumentation Review

Owner: Analytics. [pending]

Events → dashboards verified on real rows; ≥1 alert test-injected (privacy, reconciliation drift, SLA breach spike).

14. GTM Plan

Owner: Business + Product.

- In-app announcement + first-run walkthrough of the three choices.
- Vernacular one-pager + captain / floor-champion briefing — including the one hard question: “why was my accepted case deducted and his wasn’t?” (answer: every accepted case goes through the same review; the FAQ owns it).
- Support briefed; decision-tree reroute live on day one (R6).
- Sequenced rollout: Pareto cohort → broader matched cohort → all Pilots.

SECTION IV · Experiment Launch Checklist

15. Experiment Launch Checklist

Owner: Product. [pending]

Theme	Item	Owner	Status
Sign-offs	All KRD section sign-offs, incl. Security (image broker)	PM	Pending
Decisions	Y%, X, [T], leakage budget locked with data	Business + Biz-Fin	Pending
Decisions	Awareness overlay copy/treatment finalised with the designer	Design + Product	Pending
Data	debit_date live and equal to payout execution date on real rows	Analytics	Pending
Data	Reason mappings + whitelisted array delivered	Business + Product	Pending
Integrations	Kapture create/update verified two-way; SLA clock + auto-waive tested (incl. Kapture-down case)	Eng + Kapture	Pending
Integrations	LiF day-0 feed + return/credit event + dedupe verified	Payout eng + Eng	Pending

Theme	Item	Owner	Status
Ops	Agent rubric + QC live; queue set; capacity confirmed	Support	Pending
Ops	Decision-tree reroute live	Support (Nikhil)	Pending
Comms	Notification approach decided (SMS/push feasibility); templates if applicable; day-0 LiF SMS keep/drop answered (v8)	Product + Opstech	Pending
Quality	UAT complete, zero open P0/P1; reconciliation scenario passed	Product + Design + Eng	Pending
Safety	Feature flag + feed kill-switches tested	Eng	Pending
Decisions	F13 pre-payout confirmation: firm or not + channel decided	Product + Business	Pending
Integrations	Payment deep-link (F27) verified end-to-end (v8)	Payout FE team + Eng	Pending

SECTION V · Post-Launch Governance

16. Post-Launch Bugs & Hotfixes

Owner: Product & Tech. [filled post-launch]

Bug log (JIRA · severity · owner · resolved) + learnings. Incident RACI covers privacy (R2) and reconciliation drift as sev-1.

17. Experimentation Results & Learnings

Owner: Product & Analytics. [filled post-experiment]

Results vs expected · assumption register · scale / iterate / roll back call. Holdout retained; 30/90-day review.

SECTION VI · Scale-Up Strategy & Governance

18. Dashboards & Alerts (Before Scale)

GATE: model-health + queue + reconciliation dashboards verified on real data; alerts test-injected (privacy, reconciliation, SLA-breach spike).

19. Scale-Up Plan

Phased: Pareto cohort → matched broader cohort → all Pilots (attribution coverage extends in step). GATE before all-Pilots: privacy clean; reconciliation clean; SLA compliance ≥ [TBD]%; grievance within bounds; queue stable at projected volume. RED FLAG: assuming adoption because it shipped (R9).

20. Scale-Up Checklist

Experiment positive · LoDs captured · UAT zero P0/P1 · reconciliation verified at scale · queue holds · support ready · privacy verified · documentation done.

Appendix

21. Changelog

Date	Change	Owner
25–29 Jun 2026	v1–v2: initial read-only report card — two-card model (Actual/Potential), gold-read service, R-rec	PM [Shreyash]
13–21 Jul 2026	v3–v5: team-review updates — status column in gold; reason scalability + whitelist; mappings; public image links; comment resolutions	PM [Shreyash]
24 Jul 2026	v6: focus locked to Pilots; two loss types with correct sources (LiF → loss engine; SX claims → gold); central-loss-DB decision; payout automation removed/deferred	PM [Shreyash]
3 Sep 2026	v7 (this doc): model replaced — dispute-before-debit with accept (Y%-of-shipments selection), single dispute + 7-day SLA + auto-waive, cool-off; tabs Needs action / Debits / Wrong Pickups; LiF day-0 in both tabs with bank-statement credit; debit_date contract + payment pointers; Kapture two-way integration + image broker + decision-tree reroute; notifications stack; V9 design referenced. Old Potential concept and §9's prior structure retired.	PM [Shreyash]
3 Sep 2026	dispute window = 7 days from attribution for every case — seen-based clock and 14-day cap removed (F5, F26, 6b, 6j, FAQ, UAT); (2) notifications downgraded to P1 with feasibility itself open (F25); working assumption = in-app discovery via payments; day-0 LiF SMS ask marked keep/drop; R1 reworked; (3) F13 pre-payout confirmation marked open decision (firmness + channel); (4) payment deep-link moved into V1 (F27, US-07, dependency, UAT-19); push notifications → V2; V1.5 bucket dissolved; (5) dispute input = chips + text only — no voice notes ; (6) awareness overlay retained , treatment with designer; (7) deferred item reworded: payout deduction for SX-claims (today manual file upload in loss engine); (8) 6c ticket-channel map added; (9) NSM clarified — counts debited types only; wrong pickups tracked as a separate advisory rate; (10) “Delivery photo not clear” = ICUD junk (resolved); (11) open-decision register expanded in 6h (badge counting, silence display, security sign-off, BTS behaviour); timeline → “To be updated”; flywheel added to §2; PM's inline edits preserved.	PM [Shreyash]

22. FAQs

- **Does money still get cut if the Pilot does nothing?** Yes — after a **7-day window from the day the loss is attributed** (notification approach under decision; Pilots open the app for payments) (v8). Silence is a choice the Pilot was clearly offered, not a trap.
- **Why did one accepted case get deducted and another not?** Every accepted case goes through the same internal review; some are deducted, some are not. The app promises the process and tells the Pilot the result before payout — it never promises a waiver. (Confirmation mechanism before payout — under decision.)

- **What if our team misses the 7-day dispute reply?** The case is auto-waived. Our delay is never the Pilot's cost.
- **Can a Pilot dispute again after losing?** No — one dispute per loss, stated before submit. Repeated wrong disputes pause the dispute option for a while; accepting stays open.
- **Why does a lost parcel show as already deducted?** Lost-in-Field is auto-deducted the same day upstream. The app's job is the remedy: return it to the hub within 7 days and the money comes back as a credit in the next cycle.
- **Why don't the tab's date totals match the payment line?** The payment line is cumulative over the Pilot's payout period; the tab groups by date. Each item therefore points to its exact payment line — match at the item level.

23. For interested readers (optional)

Skippable detail — the doc above is complete without it.

- **Why silence = deduction is fair here.** The alternative (silence = waive) means nothing ever gets deducted and deterrence dies. Fairness is delivered by the protections instead: notice (SMS + reminders), a window that starts on sight, a hard cap, and an SLA that penalises us, not the Pilot.
- **Why Y% of shipments, not Y% of amount.** A per-amount haircut promises a discount — which the business did not want to advertise. Per-shipment selection keeps recovery at the pool level while each case stays all-or-nothing; the cost is explainability, which the FAQ and captain briefing carry.
- **Why one dispute, reviewed 100%.** The accept option absorbs genuine mistakes cheaply, so disputes self-sort toward genuinely contested cases — volumes small enough for full human review, no sampling machinery.
- **Why disputes are also a quality tool.** Overturn rate by reason and by agent is a free, labelled audit of the attribution pipeline; if one reason keeps getting overturned, the fix is upstream, not Pilot coaching.
- **Why the SLA clock must be ours.** With auto-waive as the breach rule, whoever owns the clock owns the money. A 3P queue outage must read as “waive”, not “wait”.
- **Reconciliation fine print.** debit_date = the payout “Payment for” date (not the credited date). Claims post via the adhoc-adjustment rail into Adjustments; that rail may change later, so all labels/sections are config — the invariant is line-level equality between the app and the payment.

Reviewer's Note (self-flag)

Assumption-heavy: sizing rests on the Pareto-cohort numbers [analytics, Jun '26]; full-base sizing blank. Y%, X, [T], leakage budget, SLA-compliance target, queue-capacity threshold — all [TBD] and gating the experiment. Agent capacity is asserted (12–15 existing) but incremental headroom is unmeasured.

Load-bearing integrations, none contracted yet: Kapture two-way APIs (+ per-instance customisation unknown), the image broker + security sign-off, and the LiF day-0 feed (SDE thinking through the payout-engine side-channel). Any one slipping moves the post-20-Nov target. The payment deep-link (F27) adds a fourth integration — a change on the payout team's surface — now on the V1 critical path.

Riskiest edge cases to pressure-test: SLA auto-waive under a Kapture outage (NFR-6); LiF dedupe when the day-7 record disagrees with the side-channel; the accepted-case confirmation racing the payout run (F13). **(v8)** Silent-lapse is now the top product risk: the window runs from attribution and notification feasibility is undecided — if notifications are ruled out, the 7-day window length itself should be revisited (6h).

Thin / scaffolded: 6d (designer review pending on V9), 7–8, 10 results, Sections V–VI. **Open decisions:** §6h list — the awareness overlay and the ICUD-mapping question are the two that can change scope.

#	Reason	Source	Image set on card
---	--------	--------	-------------------

1	ICUD image junk	SX claims (gold)	2 images
2	Pickup image junk	SX claims (gold)	2 images
3	Pickup-vs-QC mismatch	SX claims (gold)	4 images (2 pickup + 2 Sec-QC)
4	Wrong-RVP pickup	SX claims (gold)	none (info only)
5	Lost in Field (shipment lost by the pilot on field)	Loss Engine (shipment-loss flow)	none (reason + tip)